

E1 (a) Testing a coin is fair corresponds to hypotheses

$$H_0: \pi = 0.5 \quad \text{vs.} \quad H_a: \pi \neq 0.5$$

The count  $X$  of heads, assuming  $H_0$  holds has the  $\text{Binom}(200, 0.5)$  distribution. We use `qbinom()` to search for the 0.025-quantile

$$\text{qbinom}(0.025, 200, 0.5) \rightarrow \text{answer: } 86$$

But, in fact, with a test statistic  $X = 86$  we would get P-value

$$2 * \text{pbinom}(86, 200, 0.5) \rightarrow \text{answer: } 0.056$$

which is not below 0.05. But

$$2 * \text{pbinom}(85, 200, 0.5) \rightarrow \text{answer: } 0.0400$$

Thus  $X \leq 85$  is the lower part of the rejection region and,

by symmetry,  $X \geq 115$  is the upper part.

Now if, in actual fact,  $\pi = 0.55$ , the `dbinom(., 200, 0.55)`

is the correct pmf to consult for the value of  $\Pr(X \leq 85 \text{ or } X \geq 115)$ .

Either of these commands yield the answer:

$$\begin{aligned} \text{sum}(\text{dbinom}(c(0:85, 115:200), 200, 0.55)) &\rightarrow \text{answer: } 0.262 \\ \text{pbinom}(85, 200, 0.55) + 1 - \text{pbinom}(114, 200, 0.55) &\nearrow \end{aligned}$$

(b) The suggestion is to replace  $\text{Binom}(200, 0.55)$  by  $\text{Norm}(110, \sqrt{200(0.5)^2})$ .

Doing so, we have

$$\Pr(X \leq 85 \text{ or } X \geq 115) \doteq \text{pnorm}(85, 110, \sqrt{50}) + 1 - \text{pnorm}(115, 110, \sqrt{50}) \doteq 0.240.$$

(c) Now

$$\begin{aligned} \Pr(X \leq 85.5 \text{ or } X \geq 114.5) &\doteq \text{pnorm}(85.5, 110, \sqrt{50}) + 1 - \text{pnorm}(114.5, 110, \sqrt{50}) \\ &\doteq 0.2625 \end{aligned}$$

noticeably closer (using continuity correction) to the answer in (a) than (b) is.

4.21 We have

$$\pi^2 > \pi^2(1-\pi) = \frac{n\pi}{n} \pi(1-\pi) \geq \frac{10}{n} \pi(1-\pi) > \frac{9}{n} \pi(1-\pi).$$

Taking square roots gives 
$$\pi > 3 \sqrt{\frac{\pi(1-\pi)}{n}}.$$

4.27 (a) There are several ways to attack this one. Under the null hypothesis

$$H_0: \pi = 0.95,$$

the count  $X$  of successes in 10000 runs has a  $\text{Binom}(10000, 0.95)$  distribution. The command to get 0.025- and 0.975-quantiles is

$$q\text{binom}(c(0.025, 0.975), 10000, 0.95)$$

yielding 9457 and 9542, but

$$p\text{binom}(c(9456, 9542), 10000, 0.95)$$

vs.

$$p\text{binom}(c(9457, 9541), 10000, 0.95)$$

shows the former to be at the edges of the rejection region. That is, he should reject  $H_0$  if his coverage rate is

0.9456 or lower, or

0.9542 or higher.

This approach, however, does not take advantage of the Central Limit Theorem, which we may do since

$$np = (10000)(0.95) > 10 \quad \text{and}$$

$$n(1-p) = (10000)(0.05) > 10.$$

Under  $H_0$ ,  $\hat{\pi} \sim \text{Norm}(0.95, \sqrt{(0.95)(0.05)/10000})$ , approximately. Solving

$$\frac{\hat{\pi} - 0.95}{\sqrt{(0.95)(0.05)/10000}} \geq 1.96 \quad \Rightarrow \quad \hat{\pi} \geq 0.9543$$

$$\frac{\hat{\pi} - 0.95}{\sqrt{(0.95)(0.05)/10000}} \leq -1.96 \quad \Rightarrow \quad \hat{\pi} \leq 0.9457$$

approximately the same values as above.

4.28 (a) The sample mean is in the center of the confidence interval:

$$\bar{x} = \frac{1}{2}(11.2 + 54.7) = 32.95$$

(b) The hypothesized value for "no weight gain", 0, is outside the 95% confidence interval. That means that, under  $H_0: \mu = 0$ , with  $\alpha = 0.05$ ,  $\bar{x} = 32.95$  is in the rejection region — so the P-value is  $< 0.05$ .

4.29 This is completely wrong. The 95% CI changes with the sample, and shrinks in width as the sample size  $n$  grows. If 10 different studies offered 10 different 95% CIs, your biologist friend would be contradicting himself making such a statement about each of the CIs.

4.32. Given that  $E(\hat{\theta}^2) = \theta^2$ , we have

$$\text{Var}(\hat{\theta}) = E(\hat{\theta}^2) - E(\hat{\theta})^2$$

$$\Rightarrow E(\hat{\theta}) = \sqrt{E(\hat{\theta}^2) - \text{Var}(\hat{\theta})} = \sqrt{\theta^2 - \text{Var}(\hat{\theta})},$$

and this latter expression is  $\hat{\theta}$  only when  $\text{Var}(\hat{\theta}) = 0$  (an uninteresting case).

4.34 Assuming the second sample is independent from the first, we have

$$\text{Var}(\bar{X} - \bar{Y}) = \text{Var}(\bar{X}) + \text{Var}(\bar{Y}) = \frac{\sigma^2}{n} + \frac{\sigma^2}{n} = \frac{2}{n} \sigma^2, \text{ and}$$

$\bar{X} - \bar{Y} \sim \text{Norm}(0, \sigma \sqrt{\frac{2}{n}})$ . The lower and upper bounds of the 95% CI are positions in  $\text{Norm}(0, \sigma \sqrt{2}/\sqrt{n})$  which correspond to standardized  $Z$ -scores

$$Z = \pm \frac{1.96 \sigma / \sqrt{n}}{\sigma \sqrt{2} / \sqrt{n}} = \pm \frac{1.96}{\sqrt{2}}.$$

The probability, then, that  $\bar{Y}$  is inside the 95% CI is

$$> \text{pnorm}(1.96/\sqrt{2}) - \text{pnorm}(-1.96/\sqrt{2}),$$

which is 0.8342.

4.44 The command

```
> farrstats(~ Sepal.Width | Species, data = iris)
> qt(0.975, df = 49)
```

tell us that the critical value  $t^* = 2.0096$ , and

Species	$\bar{X}$	s	n	95% CI
setosa	3.428	0.3791	50	(3.320, 3.536)
versicolor	2.770	0.3138	50	(2.681, 2.859)
virginica	2.974	0.3225	50	(2.882, 3.066)

The final column of confidence intervals are computed as

$$\bar{X} \pm t^* \frac{s}{\sqrt{50}}.$$

Since there is no overlap between CIs, it seems reasonable to expect the (population) mean sepal width is different across each of the 3 species.